

data_cleaning_wrangling

July 9, 2025

1 Data Cleaning and Wrangling Practice

1.0.1 Objectives:

- Load, clean, and explore real-world datasets using Python and pandas.
- Generate insights through descriptive statistics and visualizations.
- Communicate findings effectively via notebooks and charts.
- Apply data analysis skills to a project using a public dataset.

1.0.2 Public dataset source:

[Kaggle Food choices and preferences of college students](#) This dataset includes information on food choices, nutrition, preferences, childhood favorites, and other information from college students. There are 126 responses from students. Data is raw and uncleaned.

```
[48]: # Importing libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_selection import mutual_info_regression
```

1.1 Discovery: Understanding the data, its structure, and what it contains

```
[16]: # Establish file path and import data
path = 'food_coded.csv'
data = pd.read_csv(path)

# Look at a snapshot of the data
data.head()
```

```
[16]:
```

	GPA	Gender	breakfast	calories_chicken	calories_day	calories_scone	\
0	2.4	2	1	430	NaN	315.0	
1	3.654	1	1	610	3.0	420.0	
2	3.3	1	1	720	4.0	420.0	
3	3.2	1	1	430	3.0	420.0	
4	3.5	1	1	720	2.0	420.0	

	coffee	comfort_food	comfort_food_reasons	\
--	--------	--------------	----------------------	---

0	1	none	we dont have comfort
1	2	chocolate, chips, ice cream	Stress, bored, anger
2	2	frozen yogurt, pizza, fast food	stress, sadness
3	2	Pizza, Mac and cheese, ice cream	Boredom
4	2	Ice cream, chocolate, chips	Stress, boredom, cravings

	comfort_food_reasons_coded	...	soup	sports	thai_food	tortilla_calories	\
0	9.0	...	1.0	1.0	1	1165.0	
1	1.0	...	1.0	1.0	2	725.0	
2	1.0	...	1.0	2.0	5	1165.0	
3	2.0	...	1.0	2.0	5	725.0	
4	1.0	...	1.0	1.0	4	940.0	

	turkey_calories	type_sports	veggies_day	vitamins	waffle_calories	\
0	345	car racing	5	1	1315	
1	690	Basketball	4	2	900	
2	500	none	5	1	900	
3	690	NaN	3	1	1315	
4	500	Softball	4	2	760	

	weight
0	187
1	155
2	I'm not answering this.
3	Not sure, 240
4	190

[5 rows x 61 columns]

```
[27]: df.describe()
```

```
[27]:
```

	GPA	Gender	grade_level	parents_cook	weight	\
count	125.000000	125.000000	125.000000	125.000000	125.000000	
mean	3.418936	1.392000	2.376000	1.528000	158.360000	
std	0.382553	0.490161	1.133536	0.746778	31.119022	
min	2.200000	1.000000	1.000000	1.000000	100.000000	
25%	3.200000	1.000000	1.000000	1.000000	135.000000	
50%	3.500000	1.000000	2.000000	1.000000	155.000000	
75%	3.700000	2.000000	3.000000	2.000000	180.000000	
max	4.000000	2.000000	4.000000	5.000000	265.000000	

	cook	eating_out	ethnic_food	fruit_day	veggies_day	\
count	125.000000	125.000000	125.000000	125.000000	125.000000	
mean	2.792000	2.560000	3.744000	4.224000	4.008000	
std	1.026236	1.13876	1.177093	0.923388	1.081337	
min	1.000000	1.000000	1.000000	1.000000	1.000000	
25%	2.000000	2.000000	3.000000	4.000000	3.000000	

50%	3.000000	2.000000	4.000000	5.000000	4.000000
75%	3.000000	3.000000	5.000000	5.000000	5.000000
max	5.000000	5.000000	5.000000	5.000000	5.000000

	healthy_feeling	self_perception_weight	vitamins	fav_cuisine_coded
count	125.000000	125.000000	125.000000	125.000000
mean	5.456000	3.120000	1.512000	2.424000
std	2.585643	1.111523	0.501867	1.947968
min	1.000000	1.000000	1.000000	0.000000
25%	3.000000	2.000000	1.000000	1.000000
50%	5.000000	3.000000	2.000000	1.000000
75%	8.000000	4.000000	2.000000	4.000000
max	10.000000	6.000000	2.000000	8.000000

```
[9]: # Preview all the columns and their quantity
print(len(data.columns))
print(data.columns)
```

61

```
Index(['GPA', 'Gender', 'breakfast', 'calories_chicken', 'calories_day',
      'calories_scone', 'coffee', 'comfort_food', 'comfort_food_reasons',
      'comfort_food_reasons_coded', 'cook', 'comfort_food_reasons_coded.1',
      'cuisine', 'diet_current', 'diet_current_coded', 'drink',
      'eating_changes', 'eating_changes_coded', 'eating_changes_coded1',
      'eating_out', 'employment', 'ethnic_food', 'exercise',
      'father_education', 'father_profession', 'fav_cuisine',
      'fav_cuisine_coded', 'fav_food', 'food_childhood', 'fries', 'fruit_day',
      'grade_level', 'greek_food', 'healthy_feeling', 'healthy_meal',
      'ideal_diet', 'ideal_diet_coded', 'income', 'indian_food',
      'italian_food', 'life_rewarding', 'marital_status',
      'meals_dinner_friend', 'mother_education', 'mother_profession',
      'nutritional_check', 'on_off_campus', 'parents_cook', 'pay_meal_out',
      'persian_food', 'self_perception_weight', 'soup', 'sports', 'thai_food',
      'tortilla_calories', 'turkey_calories', 'type_sports', 'veggies_day',
      'vitamins', 'waffle_calories', 'weight'],
      dtype='object')
```

1.2 Cleaning: Modifying data that is incorrect, incomplete, irrelevant, duplicated, or improperly formatted

```
[17]: # Convert GPA and weight data from object to floats
data['GPA'] = pd.to_numeric(data['GPA'], errors='coerce')
# print(data['GPA'])

data['weight'] = pd.to_numeric(data['weight'], errors='coerce')
# print(data['weight'])
```

```
[18]: # Identify open-ended questions and remove data entries
text_cols = data.select_dtypes(include=['object', 'string']).columns
print(text_cols)
print('Removing ',len(text_cols), ' columns')

df = data.drop(columns=text_cols)
print(df.columns)
print('Now with ',len(df.columns), ' columns')

Index(['comfort_food', 'comfort_food_reasons', 'diet_current',
      'eating_changes', 'father_profession', 'fav_cuisine', 'food_childhood',
      'healthy_meal', 'ideal_diet', 'meals_dinner_friend',
      'mother_profession', 'type_sports'],
      dtype='object')
Removing 12 columns
Index(['GPA', 'Gender', 'breakfast', 'calories_chicken', 'calories_day',
      'calories_scone', 'coffee', 'comfort_food_reasons_coded', 'cook',
      'comfort_food_reasons_coded.1', 'cuisine', 'diet_current_coded',
      'drink', 'eating_changes_coded', 'eating_changes_coded1', 'eating_out',
      'employment', 'ethnic_food', 'exercise', 'father_education',
      'fav_cuisine_coded', 'fav_food', 'fries', 'fruit_day', 'grade_level',
      'greek_food', 'healthy_feeling', 'ideal_diet_coded', 'income',
      'indian_food', 'italian_food', 'life_rewarding', 'marital_status',
      'mother_education', 'nutritional_check', 'on_off_campus',
      'parents_cook', 'pay_meal_out', 'persian_food',
      'self_perception_weight', 'soup', 'sports', 'thai_food',
      'tortilla_calories', 'turkey_calories', 'veggies_day', 'vitamins',
      'waffle_calories', 'weight'],
      dtype='object')
Now with 49 columns
```

```
[19]: # Check for null values
print(df.isna().sum())
```

GPA	5
Gender	0
breakfast	0
calories_chicken	0
calories_day	19
calories_scone	1
coffee	0
comfort_food_reasons_coded	19
cook	3
comfort_food_reasons_coded.1	0
cuisine	17
diet_current_coded	0
drink	2
eating_changes_coded	0

```

eating_changes_coded1      0
eating_out                  0
employment                  9
ethnic_food                 0
exercise                    13
father_education            1
fav_cuisine_coded          0
fav_food                    2
fries                       0
fruit_day                   0
grade_level                 0
greek_food                  0
healthy_feeling             0
ideal_diet_coded           0
income                      1
indian_food                 0
italian_food                0
life_rewarding              1
marital_status              1
mother_education            3
nutritional_check           0
on_off_campus               1
parents_cook                0
pay_meal_out                0
persian_food                1
self_perception_weight      1
soup                        1
sports                      2
thai_food                   0
tortilla_calories           1
turkey_calories             0
veggies_day                 0
vitamins                    0
waffle_calories             0
weight                      5
dtype: int64

```

```

[22]: # Decisions based on data visualization performed down below
      # GPA and weight replace with either median because data are skewed and there
      ↪are several outliers
      df['GPA'] = df['GPA'].fillna(df['GPA'].median())
      df['weight'] = df['weight'].fillna(df['weight'].median())

      # Self-perception of weight and Cook replace with most common class (mode) for
      ↪categorical
      df['self_perception_weight'] = df['self_perception_weight'].
      ↪fillna(df['self_perception_weight'].mode()[0])

```

```
df['cook'] = df['cook'].fillna(df['cook'].mode()[0])

# Remove marital status, very heavily skewed and not useful for college student,
↳ population
# Remove more columns with too many missing values
remove_NAcolumns = ['marital_status', 'exercise', 'cuisine',
↳ 'comfort_food_reasons_coded']
df = df.drop(columns=remove_NAcolumns)
```

```
[ ]: # Significantly cut down on the size of the dataset
# Keep only columns of interest
columns_of_interest =
↳ ['GPA', 'Gender', 'grade_level', 'parents_cook', 'weight', 'cook', 'eating_out', 'ethnic_food', 'fr
df = df[columns_of_interest]

# Verify no null values left
print(df.isna().sum())
```

```
GPA                0
Gender             0
grade_level        0
parents_cook        0
weight             0
cook              0
eating_out         0
ethnic_food        0
fruit_day          0
veggies_day        0
healthy_feeling    0
self_perception_weight  0
vitamins           0
fav_cuisine_coded  0
dtype: int64
```

```
[26]: # Check to see if there are any duplicate rows
df.duplicated().sum()
```

```
[26]: 0
```

1.3 Exploratory data analysis and Visualization: descriptive statistics, correlations, basic visualizations


```
[28]: # Display a few averages
ave_GPA = np.mean(df['GPA'])
ave_weight = np.mean(df['weight'])
print('Average GPA is:', ave_GPA)
print('Average weight is', ave_weight, 'lbs')
```

Average GPA is: 3.418936
Average weight is 158.36 lbs

```
[30]: # Visualize marital status
# Count occurrences
marital_counts = data['marital_status'].value_counts()

labels = ['Single', 'In a_
↪relationship', 'Cohabiting', 'Married', 'Divorced', 'Widowed']

# Plot bar chart
plt.figure(figsize=(8, 5))
plt.bar(marital_counts.index, marital_counts.values, color='skyblue')
plt.title('Marital Status Distribution')
plt.xlabel('Marital Status')
plt.ylabel('Count')
plt.tight_layout()
plt.xticks(ticks=range(1, 7), labels=labels)
plt.show()
```



```
[32]: # Visualize Gender
# Count occurrences
gender_counts = data['Gender'].value_counts()
# The index is the unique categories
# The values are the number of occurrences
```

```

# Plot pie chart
fig, ax = plt.subplots()
ax.pie(gender_counts.values, labels=['Female', 'Male'], autopct='%1.1f%%',
      colors=['lightpink', 'thistle'])
ax.set_title('Gender Distribution')
plt.show()

# Count occurrences
grade_counts = data['grade_level'].value_counts()
# The index is the unique categories
# The values are the number of occurrences


labels = ['Freshman', 'Sophomore', 'Junior', 'Senior']

# Plot bar chart
plt.figure(figsize=(8,5))
bars = plt.bar(grade_counts.index, grade_counts.values, color='lightsalmon')
total = grade_counts.values.sum()
for bar in bars:
    height = bar.get_height()
    percentage = 100 * height / total
    plt.text(bar.get_x() + bar.get_width()/2, height + 1, f'{percentage:.1f}%',
             ha='center', va='bottom')

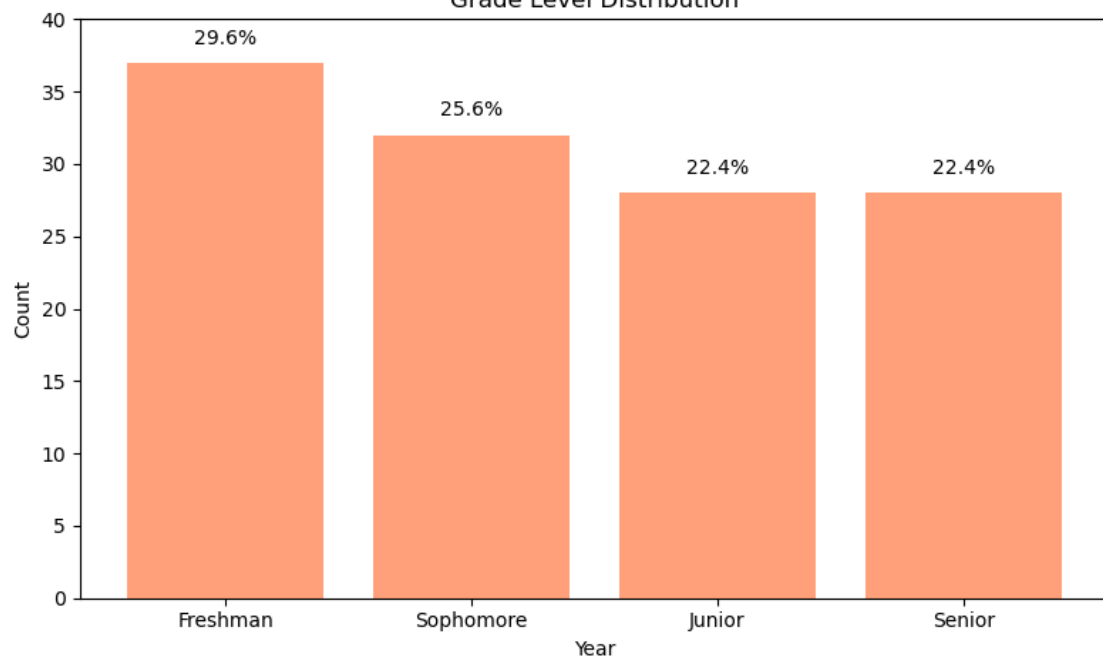
plt.title('Grade Level Distribution')
plt.xlabel('Year')
plt.ylabel('Count')
plt.tight_layout()
plt.ylim(0,40)
plt.xticks(ticks=range(1, 5), labels=labels)
plt.show()

```



Gender Distribution



Grade Level Distribution





```
[33]: # Visualize GPA
x = data['GPA']
y = data['weight']
plt.scatter(x, y, color = 'mediumaquamarine')
plt.xlabel('GPA')
plt.ylabel('Weight')
plt.title('Weight vs GPA')
plt.show()
```



```
[34]: # Visualize distributions of weight and GPA
sns.histplot(data['weight'], bins=30, kde=True, stat="density",
             color='skyblue', edgecolor='darkblue')
plt.title('Weight Distribution')
plt.show()

sns.histplot(data['GPA'], bins=15, kde=True, stat="density", color='lightpink',
             edgecolor='firebrick')
plt.title('GPA Distribution')
plt.show()
```





```
[37]: # Visualizing cooking variables
parents_counts = data['parents_cook'].value_counts()
labels = ['Almost everyday', '2-3x/week', '1-2x/week', 'Holidays only', 'Never']
plt.figure(figsize=(8,5))
bars1= plt.bar(parents_counts.index, parents_counts.values, color='teal')
total = parents_counts.values.sum()
for bar in bars1:
    height = bar.get_height()
    percentage = 100 * height / total
    plt.text(bar.get_x() + bar.get_width()/2, height + 1, f'{percentage:.1f}%',
             ha='center', va='bottom')

plt.title('Approximately how many days a week did parents cook?')
plt.xlabel('Frequency')
plt.ylabel('Count')
plt.ylim(0,80)
plt.tight_layout()
plt.xticks(ticks=range(1, 6), labels=labels)
plt.show()
```

```

cook_counts = data['cook'].value_counts()
labels = ['Everyday', 'Couple times/week', 'Infrequently', 'Holidays only', 'Never']
plt.figure(figsize=(8,5))
bars2 = plt.bar(cook_counts.index, cook_counts.values, color='cadetblue')
total = cook_counts.values.sum()
for bar in bars2:
    height = bar.get_height()
    percentage = 100 * height / total
    plt.text(bar.get_x() + bar.get_width()/2, height + 1, f'{percentage:.1f}%',
             ha='center', va='bottom')



plt.title('How often do you cook?')
plt.xlabel('Frequency')
plt.ylabel('Count')
plt.tight_layout()
plt.ylim(0,55)
plt.xticks(ticks=range(1, 6), labels=labels)
plt.show()

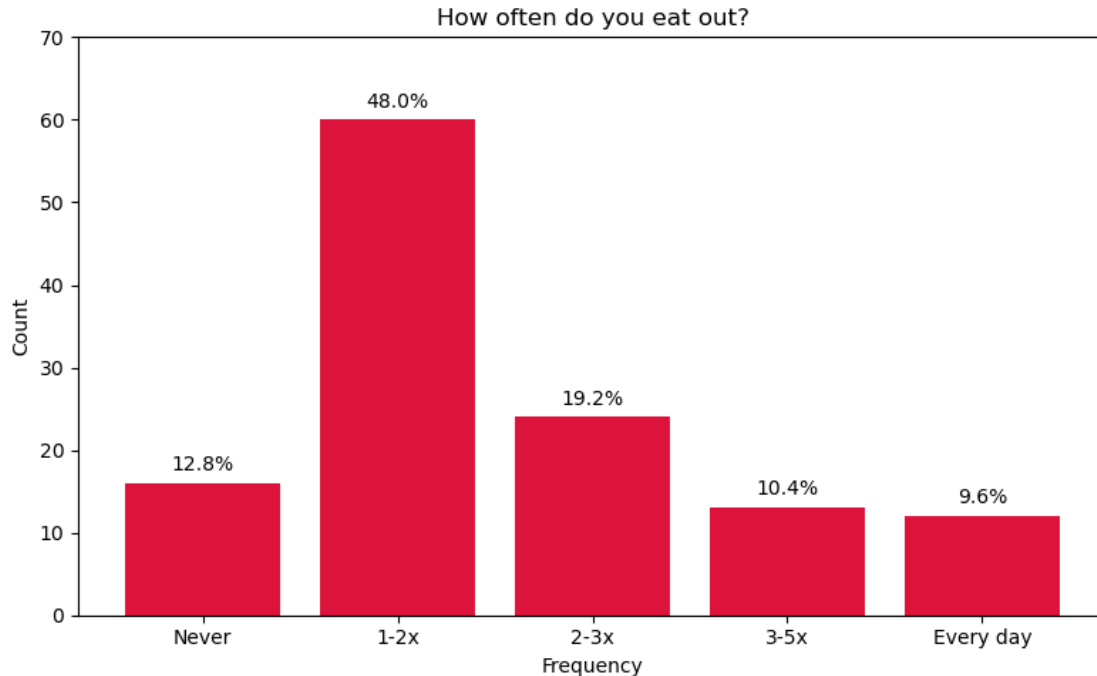
eatingout_counts = data['eating_out'].value_counts()
labels = ['Never', '1-2x', '2-3x', '3-5x', 'Every day']

plt.figure(figsize=(8,5))
bars = plt.bar(eatingout_counts.index, eatingout_counts.values, color='crimson')
# Add percentage labels
total = eatingout_counts.values.sum()
for bar in bars:
    height = bar.get_height()
    percentage = 100 * height / total
    plt.text(bar.get_x() + bar.get_width()/2, height + 1, f'{percentage:.1f}%',
             ha='center', va='bottom')

# plt.bar(eatingout_counts.index, eatingout_counts.values, color='crimson')
plt.title('How often do you eat out?')
plt.xlabel('Frequency')
plt.ylabel('Count')
plt.ylim(0,70)
plt.tight_layout()
plt.xticks(ticks=range(1, 6), labels=labels)
plt.show()

```





```
[38]: # Visualize fruit consumption
labels=['Very unlikely','Unlikely','Neutral','Likely','Very Likely']
fruit_counts = data['fruit_day'].value_counts().sort_index()
print(fruit_counts)

colors = ['#e6194b', # red
          '#f58231', # orange
          '#ffe119', # yellow
          '#3cb44b', # green
          '#4363d8'] # blue


fig, ax = plt.subplots()
ax.pie(fruit_counts.values, labels=labels, autopct='%1.1f%%', colors=colors)
ax.set_title("How likely are you to eat fruit in a day?")
plt.show()

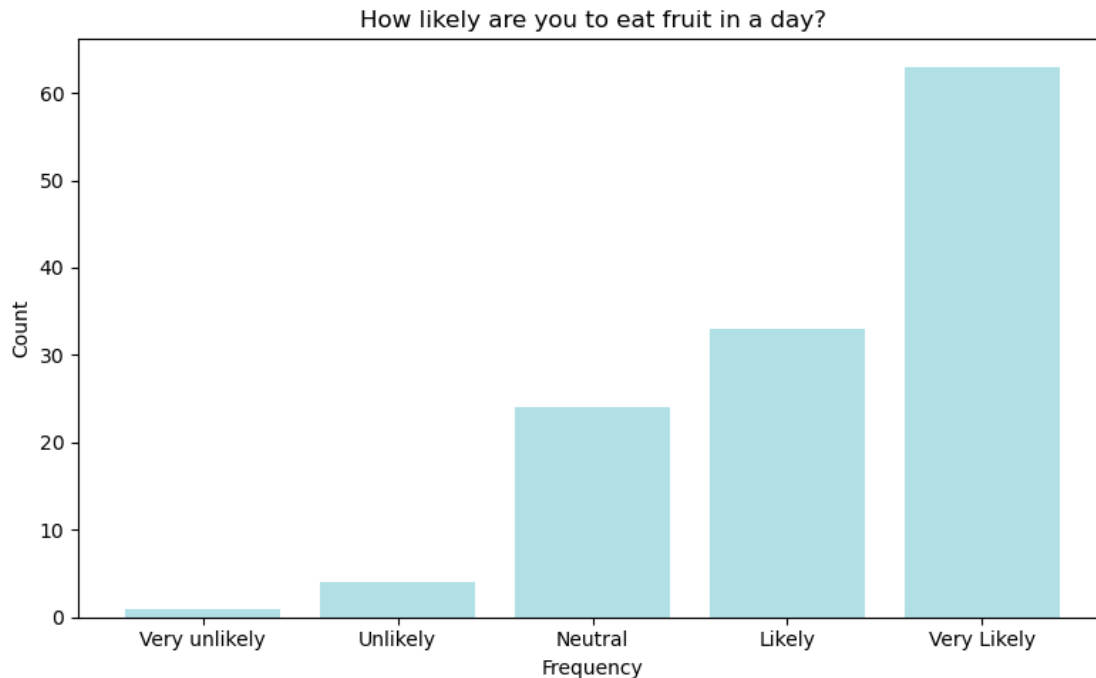
plt.figure(figsize=(8,5))
plt.bar(fruit_counts.index, fruit_counts.values, color='powderblue')
plt.title('How likely are you to eat fruit in a day?')
plt.xlabel('Frequency')
plt.ylabel('Count')
plt.tight_layout()
plt.xticks(ticks=range(1, 6), labels=labels)
plt.show()
```

fruit_day

```
1      1
2      4
3     24
4     33
5     63
Name: count, dtype: int64
```

How likely are you to eat fruit in a day?





```
[39]: # Visualize vegetable consumption
labels=['Very unlikely','Unlikely','Neutral','Likely','Very Likely']
veggies_counts = data['veggies_day'].value_counts().sort_index()
print(fruit_counts)

colors = ['lightcoral', # red
          'wheat', # orange
          'lemonchiffon', # yellow
          'mediumspringgreen', # green
          'lightskyblue'] # blue


fig, ax = plt.subplots()
ax.pie(veggies_counts.values, labels=labels, autopct='%1.1f%%', colors=colors)
ax.set_title("How likely are you to eat vegetables in a day?")
plt.show()


plt.figure(figsize=(8,5))
plt.bar(veggies_counts.index, veggies_counts.values, color='lightpink')
plt.title('How likely are you to eat vegetables in a day?')
plt.xlabel('Frequency')
plt.ylabel('Count')
plt.tight_layout()
plt.xticks(ticks=range(1, 6), labels=labels)
plt.show()
```

fruit_day

```
1      1
2      4
3     24
4     33
5     63
Name: count, dtype: int64
```


How likely are you to eat vegetables in a day?





```
[40]: # Visualize self-perception of weight
perception_counts = data['self_perception_weight'].value_counts().sort_index()

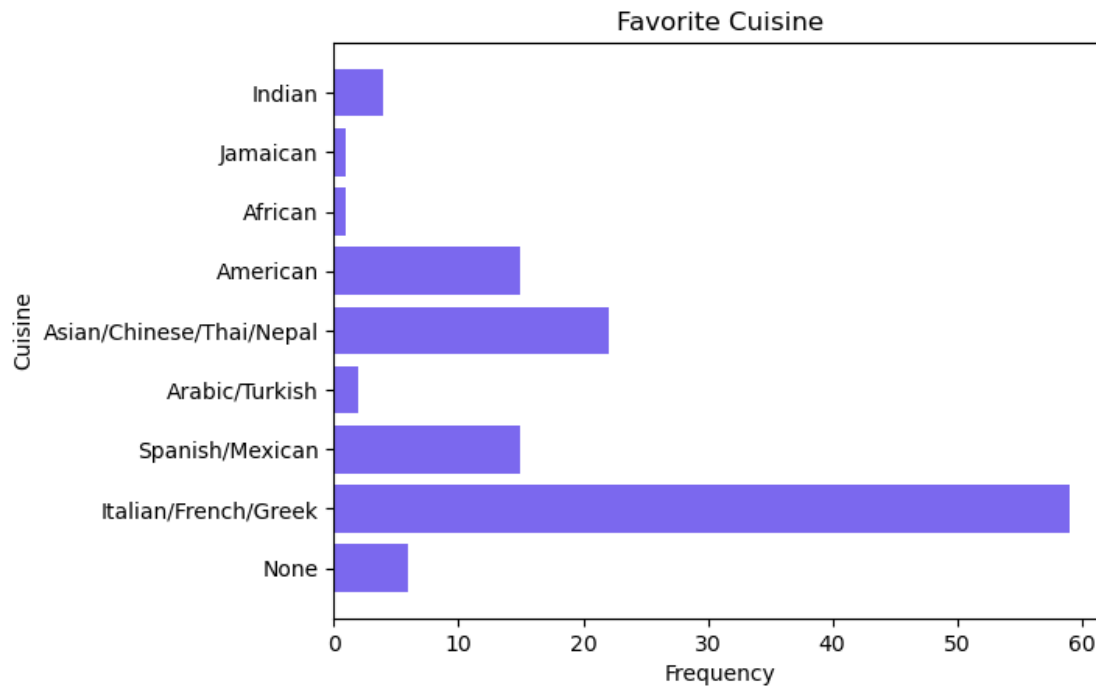
labels = ['Slim', 'Very fit', 'Just right', 'Slightly_
↳ overweight', 'Overweight', 'None']
# 'None' corresponds to: "I don't think of myself in these terms"
plt.scatter(df['self_perception_weight'], df['weight'])
plt.xticks(ticks=range(1, 7), labels=labels, rotation=45)
plt.title('Self perception of weight')
plt.ylabel('Frequency')
plt.show()
```




```
[42]: # Visualize vitamins and favorite cuisine
vitamins_counts = data['vitamins'].value_counts().sort_index()
print(vitamins_counts)

favcuisine_counts = data['fav_cuisine_coded'].value_counts().sort_index()
labels = ['None', 'Italian/French/Greek', 'Spanish/Mexican', 'Arabic/
↳Turkish', 'Asian/Chinese/Thai/Nepal', 'American', 'African', 'Jamaican', 'Indian']
# 'None' corresponds to: "I don't think of myself in these terms"
plt.barh(favcuisine_counts.index, favcuisine_counts.values,
↳color='mediumslateblue')
plt.yticks(ticks=range(0, 9), labels=labels)
plt.title('Favorite Cuisine')
plt.ylabel('Cuisine')
plt.xlabel('Frequency')
plt.show()
```

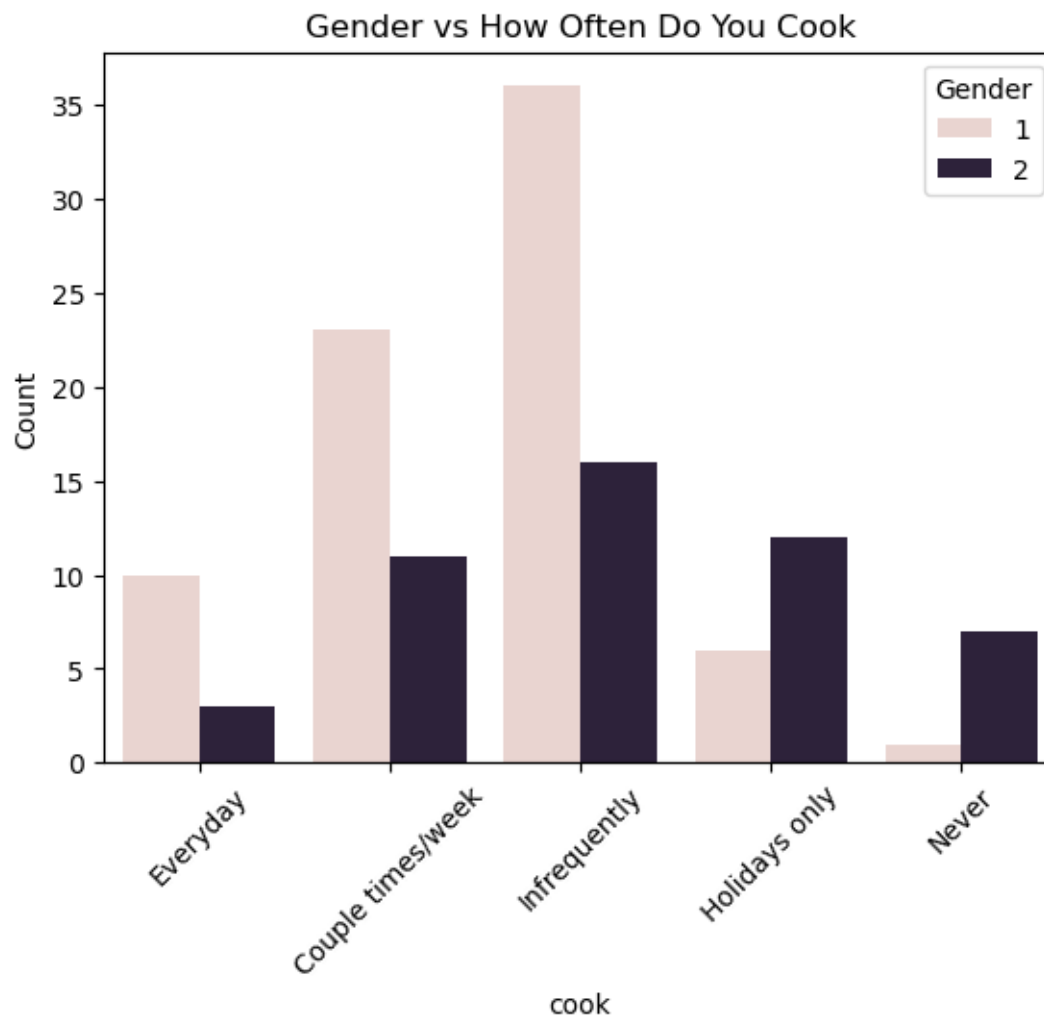
```
vitamins
1    61
2    64
Name: count, dtype: int64
```



```
[44]: # Visualize gender and weight
sns.boxplot(x=df['Gender'], y=df['weight'], color='mediumaquamarine')
plt.xlabel('Gender')
plt.ylabel('Weight')
plt.title('Gender vs Weight')
plt.xticks(ticks=[0, 1], labels=['Female', 'Male'])
plt.show()
```




```
[45]: # Visualize gender and cooking
sns.countplot(x='cook', hue='Gender', data=df)
plt.title('Gender vs How Often Do You Cook')
plt.ylabel('Count')
plt.xticks(ticks=range(0,5), labels = ['Everyday', 'Couple times/
↪ week', 'Infrequently', 'Holidays only', 'Never'], rotation=45)
plt.show()
```




```
[46]: #Pearson correlation
corr_matrix = df.corr(method='pearson')

plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f", square=True)
plt.title("Pearson Correlation Matrix")
plt.show()
```



```
[47]: # Create a mask for correlations with abs value < 0.3
mask = (np.abs(corr_matrix) < 0.3)

# Plot
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix,
            mask=mask,
            annot=True,
            cmap='RdBu_r',
            fmt=".2f",
            square=True,
            linewidths=0.5,
            cbar=True)
plt.title("Filtered Pearson Correlation Matrix (|r| < 0.3)")
plt.show()
```

```
[49]: X = df.copy() # make a copy
X = X.dropna() # remove rows with NA for now
X = X.apply(lambda col: col.astype(int) if col.dtype == 'float' else col) #
    ↳convert all floats to ints
y = X.pop("GPA")

# All discrete features should now have integer dtypes (double-check this
    ↳before using MI!)
discrete_features = X.dtypes == int

def make_mi_scores(X, y, discrete_features):
    mi_scores = mutual_info_regression(X, y,
    ↳discrete_features=discrete_features)
    mi_scores = pd.Series(mi_scores, name="MI Scores", index=X.columns)
    mi_scores = mi_scores.sort_values(ascending=False)
```

```

    return mi_scores

mi_scores = make_mi_scores(X, y, discrete_features)
print(mi_scores)

def plot_mi_scores(scores):
    scores = scores.sort_values(ascending=True)
    width = np.arange(len(scores))
    ticks = list(scores.index)
    plt.barh(width, scores)
    plt.yticks(width, ticks)
    plt.title("Mutual Information Scores")


plt.figure(dpi=100, figsize=(8, 5))
plot_mi_scores(mi_scores)

```

```

weight                1.356273
self_perception_weight 0.176837
fav_cuisine_coded     0.163722
grade_level           0.118304
ethnic_food           0.112594
eating_out            0.095915
healthy_feeling       0.048478
cook                  0.043687
Gender                0.041797
vitamins              0.006173
parents_cook          0.000000
fruit_day             0.000000
veggies_day           0.000000
Name: MI Scores, dtype: float64

```



Interpretation: Mutual Information quantifies how much information one variable gives about another.

A higher MI score corresponds to a stronger non-linear and linear dependency with the target.

According to the plot above showing how much each feature contributes to predicting **GPA** using mutual information.

Most Informative Features:

* weight (MI 1.71): most informative variable for predicting the target * self_perception_weight (MI 0.32): moderately informative, reflects how someone's perception of their body relates to the target. healthy_feeling (MI 0.13): provides some predictive signal, although weaker.

Less Informative Features:

* veggies_day, cook, fruit_day, ethnic_food: have low MI values (< 0.1), suggesting a weak relationship with the target.

Non-informative Features: * Gender, grade_level, parents_cook, eating_out, vitamins (MI = 0): these variables do not contribute any useful information about the target

1.4 Summarized notes, insights, and interpretation

- Most common favorite cuisine is Italian/French/Greek
- Gender and cooking were highly correlated ($r=0.31$), as was gender and weight (as expected) ($r=0.51$).
- Ethnic food and veggies day ($r=0.35$) and fruit day ($r=0.31$) are highly correlated
- Fruit day and veggies day are highly correlated ($r=0.65$)
- The distribution of population is 60.8% female, 39.2% male
- Weight distribution is skewed to the right
- GPA distribution is skewed to the left

- Grade level distribution is somewhat evenly distributed (29.6% freshman, 25.6% sophomores, 22.4% juniors, and 22.4% seniors)
- 60% of students reported that their parents cooked almost everyday of the week (parents_cook)
- 10% of students reported that they cook every day (cook)
- 42% of students reported that they cook infrequently
- 48% of students reported that they eat out 1-2 times a week (eating_out)
- 10% of students reported that they eat out every day